

Introduction to Semantic Analysis

KC Sivaramakrishnan

IIT Madras

Monsoon 2026

The Compiler so far

- Lexical Analysis
 - Detects inputs with illegal tokens
- Parsing
 - Detects inputs with ill-formed parse trees
- **Next:** Semantic Analysis
 - Last 'front-end' phase
 - Catches all the remaining errors

The compilation process is driven by the syntactic structure of the program as discovered by the parser

Semantic routines:

- interpret (partial) meaning of the program based on its syntactic structure
- two purposes:
 - finish analysis by deriving context-sensitive information (e.g. type checking)
 - begin synthesis by generating the IR or target code

What does Semantic Analysis do?

- Checks of many kinds:
 - All identifiers are declared.
 - Type checking in assignment statements, function calls, expressions, etc.
 - Inheritance relationships.
 - Classes defined only once.
 - Methods in a class defined only once.
 - Reserved identifiers are not misused

And many others...

Why a Separate Semantic Analysis?

Parsing cannot catch some errors. Some language constructs are not context-free.

- All identifiers are declared before use.
 - Corresponds to the abstract language $\{w c w \mid w, c \in \{a, b\}^*\}$.
- The number of formal parameters in the declaration of a function agrees with the number of actual parameters in the use of the function.
 - Corresponds to the abstract language $\{a^n b^m c^n d^m\}$.
 - Here, a^n, b^m represent the formal-parameter list of two functions, while c^n, d^m represent the actual-parameter list in calls to the functions.

These languages are not context-free.

Scope

- Matching identifier declarations with uses
 - One of the most important tasks performed by Semantic Analysis.
- The scope of an identifier is the portion of a program in which that identifier is accessible.
- The same identifier may refer to different things in different parts of the program.
 - Different scopes for same name don't overlap.
- An identifier may have restricted scope even if declared before use.

Static vs. Dynamic Scope

- Most languages have static scope.
 - Scope depends only on the program text, not run-time behavior.
 - C,C++,Java, etc. have static scope.

Example:

```
int x = 0;
{
    y = x;
    int x = 1;
    z = x;
}
```

Uses of `x` refer to closest enclosing definition.

Static vs. Dynamic Scope

- Some languages have dynamic scope.
 - Scope depends on the execution of the program.
- A dynamically scoped variable refers to the closest enclosing binding in the execution of the program.
 - Lisp, SNOBOL have dynamic scoping for variables.
- The exception handler for an exception is dynamically determined.
- Object-oriented languages have limited dynamic scoping for identifying method bindings.

Example of a language with dynamic scoping for variables:

```
a ← 0
g(y) = let a ← 4 in f();
h(z) = let a ← 5 in f();
f() = a;
```


Symbol tables

For compile-time efficiency, compilers use a symbol table:

- associates lexical names (symbols) with their attributes

What items belong to the symbol table?

- variable names
- defined constants
- procedure and function names
- source text labels
- compiler-generated temporaries

(we'll get there)

A symbol table is a compile-time structure

Symbol table information

What kind of information might the compiler need?

- textual name
- data type
- declaring procedure
- lexical level of declaration
- storage class
- offset in storage
- can it be aliased? to what other names?
- number and type of arguments to functions
- ...

(base address)

Storage classes of variables

Variables in the program may be assigned different storage classes based on their lifetime, scope, where they live, etc.

Common storage classes are: global, stack, static, registers.

The storage classes also determine how the variables are addressed (addressing mode).

- A local variable is not assigned a fixed machine address (or relative to the base of a module)
- Rather a stack location that is accessed by an offset from a register whose value may be different every time the procedure is invoked.

Symbol table organization

How should the table be organized?

- Linear List

- $O(n)$ probes per lookup
- easy to expand — no fixed size
- one allocation per insertion

- Ordered Linear List

- $O(\log_2 n)$ probes per lookup using binary search
- insertion is expensive (to reorganize list)

- Binary Tree

- $O(n)$ probes per lookup — unbalanced
- $O(\log_2 n)$ probes per lookup — balanced
- easy to expand — no fixed size
- one allocation per insertion

- Hash Table

- $O(1)$ probes per lookup — on average
- expansion costs vary with specific scheme

Nested scopes: block-structured symbol tables

What information is needed?

- when asking about a name, want most recent declaration
- declaration may be from current scope or outer scope
- innermost scope overrides outer scope declarations

Key point: new declarations occur only in current scope

Nested scopes: block-structured symbol tables

What operations do we need?

- `void put (Symbol key, Object value)`
bind key to value
- `Object get (Symbol key)`
return value bound to key
- `void beginScope()`
remember current state of table
- `void endScope()`
close current scope and restore table to state at most recent open `beginScope`

Attributes are internal representation of declarations

Symbol table associates names with attributes

Names may have different attributes depending on their meaning:

- variables: type, procedure level, frame offset
- types: type descriptor, data size/alignment
- constants: type, value
- procedures: formals (names/types), result type, block information (local decls.), frame size

Types

- What is a type?
 - A set of values.
 - A set of operations on those values.
- Classes (e.g. in Java) are one instantiation of the modern notion of type.

Type system

Type systems is a logical system comprising of a set of rules that assigns a property called a type to every term in the language.

Why do we need Type Systems?

- Not all operations are legal for all types.
- It makes sense to add two integers, but it doesn't make any sense to add a function pointer and an integer in C.
- But both have the same assembly language implementation:
`add $r1, $r2, $r3.`

A language's type system specifies which operations are valid for which types.

- The goal of *Type Checking* is to ensure that operations are used with correct types.

Type Checking Overview

- Three kinds of languages:
 - *Statically typed*: All or almost all checking of types is done as part of compilation (C, Java, OCaml, Rust, MiniJava).
 - *Dynamically typed*: Almost all checking of types is done as part of program execution (Python, JavaScript, Scheme)
 - *Untyped*: No type checking (machine code)

Static VS Dynamic Type Checking

- Static checking catches many programming errors at compile time.
- Static checking also avoids overhead of runtime type checks.
- Dynamic type systems are more flexible than static type systems
 - allow more programs to be executed.
- Dynamically typed languages allow rapid prototyping, but may be difficult to develop and maintain large code bases.
- Recent trend: bolt-on static type checking on top of dynamically typed languages.
 - See Hack (Meta, PHP), Flow (Meta, JavaScript), TypeScript (Microsoft, JavaScript), Eqvalizer (WhatsApp, Erlang), Pyre (Meta, Python).

This is a type war. There is no clear winner.

Type Checking and Type Inference

- *Type Checking* is the process of verifying fully typed programs.
- *Type Inference* is the process of filling in missing type information.
- The two are different, but the terms are often (incorrectly) used interchangeably.
- The formalism to express both type-checking and type-inference is *logical rules of inference*.
 - Just like how the formal notation for lexical analysis and parsing were regular expressions and context-free grammars respectively.

Rules of Inference

- Inference rules have the form:
If Hypothesis is true, then Conclusion is true.
 - Hypothesis is also called as *antecedent*
 - Conclusion is also called as *consequent*
- We use the notation $\vdash e : T$ to denote that e has type T , and the notation $\vdash e$ to denote that e type-checks.
- Inference rules are often written as follows:
$$\frac{\vdash \text{Hypothesis}_1 \quad \vdash \text{Hypothesis}_2 \quad \dots \quad \vdash \text{Hypothesis}_n}{\vdash \text{Conclusion}}$$

This translates to: **If** Hypothesis₁ **and** Hypothesis₂ **and** ... Hypothesis_n **then** Conclusion.

Inference rules are a very succinct way to precisely state type-checking rules.

Rules of Inference

$$\frac{\vdash E1 : T1 \quad \vdash E2 : T2}{\vdash E3 : T3}$$

- For type inference, the rule is read as:
 - *If $E1$ has type $T1$ and $E2$ has type $T2$ then $E3$ has type $T3$.*
- For type checking, the rule is read as:
 - *If E_1 with type $T1$ and E_2 with type $T2$ type-check, then E_3 also type-checks with type $T3$.*
 - We may sometimes omit $T3$ when type checking (statements, for example).

Examples

$$\frac{\text{program contains } \text{int } i}{\vdash i : \text{int}}$$
$$\frac{\vdash e_1 : \text{int} \quad \vdash e_2 : \text{int}}{\vdash e_1 + e_2 : \text{int}}$$
$$\frac{\vdash id : \text{int} \quad \vdash e : \text{int}}{\vdash id = e}$$

Type Environment

- A type environment is a mapping from identifiers to types.
 - We will use the symbol A to denote type environment.
 - $dom(A)$ denotes the domain of A , i.e. the set of identifiers for which types are defined in the environment A .
 - We use the notation $A \vdash id : T$ to denote that $A(id) = T$.

Type expressions

Type expressions are a textual representation for types:

- ① basic types: *boolean*, *char*, *integer*, *real*, etc.
- ② type names
- ③ constructed types (constructors applied to type expressions):
 - ① *array*(I, T) denotes an array of T indexed over I
e.g., *array*($1 \dots 10, \text{integer}$)
 - ② products: $T_1 \times T_2$ denotes Cartesian product of type expressions T_1 and T_2
 - ③ records: fields have names
e.g., *record*(($a \times \text{integer}$), ($b \times \text{real}$))
 - ④ pointers: *pointer*(T) denotes the type “pointer to an object of type T ”
 - ⑤ functions: $D \rightarrow R$ denotes the type of a function mapping domain type D to range type R
e.g., $\text{integer} \times \text{integer} \rightarrow \text{integer}$

In Minijava, the type expressions are `boolean`, `int`, `int[]`, `id`.

MiniJava Type Judgements

$\vdash g$ The goal g type checks

$\vdash mc$ The main class mc type checks

$\vdash d$ The type declaration d type checks

$C \vdash m$ If defined under class C , the method m type checks

$A, C \vdash s$ In a type environment A , if written in class C ,
the statement s type checks

$A, C \vdash e : t$ In a type environment A , if written in class C ,
the expression e has type t

$A, C \vdash p : t$ In a type environment A , if written in class C ,
the primary expression e has type t

Typechecking a Minijava program

A program in Minijava consists of a main class (which contains the `main` method) and a bunch of other classes.

Our goal is to typecheck all the classes:

$$\frac{\text{distinct}(\text{classname}(mc), \text{classname}(d_1), \dots, \text{classname}(d_n)) \quad \vdash mc \quad \vdash d_i, i \in \{1, \dots, n\}}{\vdash mc \ d_1 \dots d_n}$$

This says that if all the class-names are distinct, and if all the classes individually type-check, then the entire program also type-checks.

Ref: The MiniJava Type System, Jens Palsberg.

Typechecking a Minijava class

$$\frac{\begin{array}{c} \text{distinct}(id_1, \dots, id_f) \\ \text{distinct}(\text{methodname}(m_1), \dots, \text{methodname}(m_k)) \\ id \vdash m_i \quad i \in \{1, \dots, k\} \end{array}}{\vdash \text{class } id \{ t_1 \text{ } id_1; \dots; t_f \text{ } id_f; m_1 \dots m_k \}}$$

$id \vdash m_i$ indicates that method m_i type-checks under a type-environment consisting of fields of class id .

Question: Does this allow a method and a field to have the same name?

Ref: The MiniJava Type System, Jens Palsberg.

“fields” helper function

$$\frac{\text{class } id \{ t_1 \ id_1; \dots; t_f \ id_f; m_1 \ \dots \ m_k \} \text{ is in the program}}{fields(id) = [id_1 : t_1, \dots, id_f : t_f]}$$

$$\frac{\text{class } id \text{ extends } id^P \{ t_1 \ id_1; \dots; t_f \ id_f; m_1 \ \dots \ m_k \} \text{ is in the program}}{fields(id) = fields(id^P) \cdot [id_1 : t_1, \dots, id_f : t_f]}$$

Typechecking a Minijava method

$$\frac{\begin{array}{c} \text{distinct}(id_1^F, \dots, id_n^F) \\ \text{distinct}(id_1, \dots, id_r) \\ A = \text{fields}(C) \cdot [id_1^F : t_1^F, \dots, id_n^F : t_n^F] \cdot [id_1 : t_1, \dots, id_r : t_r] \\ A, C \vdash s_i, i \in \{1, \dots, q\} \quad A, C \vdash e : t \end{array}}{C \vdash \text{public } t \text{ id}^M (t_1^F \text{ id}_1^F, \dots, t_n^F \text{ id}_n^F) \{ \\ t_1 \text{ id}_1; \dots; t_r \text{ id}_r; s_1 \dots s_q; \text{return } e; \}}$$

Question:

- Does this allow a formal parameter and a local variable to have the same name?
 - What will be its type?
- Does this allow a local variable and a field to have the same name?

Ref: The MiniJava Type System, Jens Palsberg.

Inheritance in Minijava

To express inheritance among classes, we will define the *subtype* relation, denoted by $\leq$.

$t_1 \leq t_2$ iff t_1 is a sub-class of t_2 .

The following rules describe the *subtype* relation:

$$\begin{array}{c} t \leq t \\[1em] \frac{t_1 \leq t_2 \quad t_2 \leq t_3}{t_1 \leq t_3} \\[1em] \frac{\text{class } C \text{ extends } D \quad \text{is in the program}}{C \leq D} \end{array}$$

Typechecking statements in Minijava

$$\frac{A(id) = t_1 \quad A, C \vdash e : t_2 \quad t_2 \leq t_1}{A, C \vdash id = e;}$$

$$\frac{A(id) = \text{int}[] \quad A, C \vdash e_1 : \text{int} \quad A, C \vdash e_2 : \text{int}}{A, C \vdash id [e_1] = e_2;}$$

$$\frac{A, C \vdash e : \text{boolean} \quad A, C \vdash s_1 \quad A, C \vdash s_2}{A, C \vdash \text{if } (e) s_1 \text{ else } s_2}$$

$$\frac{A, C \vdash e : \text{boolean} \quad A, C \vdash s}{A, C \vdash \text{while } (e) s}$$

$$\frac{A, C \vdash e : \text{int}}{A, C \vdash \text{System.out.println}(e);}$$

Ref: The MiniJava Type System, Jens Palsberg.

Typechecking expressions in Minijava

$$\frac{A, C \vdash p_1 : \text{boolean} \quad A, C \vdash p_2 : \text{boolean}}{A, C \vdash p_1 \ \&\& \ p_2 : \text{boolean}}$$

$$\frac{A, C \vdash p_1 : \text{int} \quad A, C \vdash p_2 : \text{int}}{A, C \vdash p_1 < p_2 : \text{boolean}}$$

$$\frac{A, C \vdash p_1 : \text{int} \quad A, C \vdash p_2 : \text{int}}{A, C \vdash p_1 + p_2 : \text{int}}$$

$$\frac{A, C \vdash p_1 : \text{int} \quad A, C \vdash p_2 : \text{int}}{A, C \vdash p_1 - p_2 : \text{int}}$$

$$\frac{A, C \vdash p_1 : \text{int} \quad A, C \vdash p_2 : \text{int}}{A, C \vdash p_1 * p_2 : \text{int}}$$

$$\frac{A, C \vdash p_1 : \text{int}[] \quad A, C \vdash p_2 : \text{int}}{A, C \vdash p_1 [p_2] : \text{int}}$$

$$\frac{A, C \vdash p : \text{int}[]}{A, C \vdash p . \text{length} : \text{int}}$$

Typechecking method calls in Minijava

$$\frac{???}{A, C \vdash p.id(e_1, e_2, \dots, e_n) : t}$$

Typechecking method calls in Minijava

$$\frac{A, C \vdash p : D \quad \text{methodtype}(D, id) = (t'_1, \dots, t'_n) \rightarrow t \quad A, C \vdash e_i : t_i \quad t_i \leq t'_i, i \in \{1, \dots, n\}}{A, C \vdash p.id(e_1, e_2, \dots, e_n) : t}$$

methodtype(*D*, *id*) returns the type of the method *id* in class *D* or one of its super-classes.

Question: What change should we make in the rule if we want to enforce access modifiers?

Typechecking method calls in Minijava with access modifiers

$$\begin{array}{c} A, C \vdash p : D \quad \text{methodtype}(D, id) = (t'_1, \dots, t'_n) \rightarrow t \\ \quad (\text{methodaccess}(D, id) = \text{public}) \\ \vee \text{methodaccess}(D, id) = \text{protected} \wedge C \leq D \\ \vee \text{methodaccess}(D, id) = \text{private} \wedge C = D \\ \hline A, C \vdash e_i : t_i \quad t_i \leq t'_i \quad i \in \{1, \dots, n\} \\ \hline A, C \vdash p.id(e_1, e_2, \dots, e_n) : t \end{array}$$

methodtype(*D*, *id*) returns the type of the method *id* in class *D* or one of its super-classes.

methodaccess(*D*, *id*) returns the access modifier of the method *id* in class *D* or one of its super-classes.

Typechecking ternary operation in Minijava

$$\frac{???}{A, C \vdash (e_1 ? e_2 : e_3) : t}$$

Typechecking ternary operation in Minijava

$$\frac{A, C \vdash e_1 : \text{boolean} \quad A, C \vdash e_2 : t \quad A, C \vdash e_3 : t}{A, C \vdash (e_1 ? e_2 : e_3) : t}$$

But what about inheritance? How to handle subtyping?

Example

```
...  
class B extends A {...}  
class C extends B {...}  
class D extends B {...}  
...  
A objA; B objB; C objC; D objD;  
...  
objA = (...) ? objC : objD;  
objB = (...) ? objC : objD;
```

Both the instances of ternary statement should type-check. This requires the ternary statement itself to have the type `B`.

Least Upper Bounds

- $\text{lub}(C, D)$, the least upper bound of classes C and D is defined to class E such that
 - ① $C \leq E$ and $D \leq E$: E is an upper bound of both C and D
 - ② $\forall E'. C \leq E' \wedge D \leq E' \Rightarrow E \leq E'$: E is least among upper bounds.
- In Minijava, the least upper bound of two classes is their least common ancestor in the inheritance tree.

Typechecking ternary operation in Minijava: Attempt-2

$$\frac{A, C \vdash e_1 : \text{boolean} \quad A, C \vdash e_2 : t_1 \quad A, C \vdash e_3 : t_2 \quad \exists t'. t_1 \leq t' \wedge t_2 \leq t' \quad t = \text{lub}(t_1, t_2)}{A, C \vdash (e_1 ? e_2 : e_3) : t}$$

Typechecking No-overloading in Minijava

- **Overloading:** A method is said to be overloaded if there exists another method in the same class or its super classes with the same name but different type.
- **Overriding:** A method in a class is said to be overridden if there exists a method in its subclass with the same name and type.

$$\frac{\begin{array}{c} \text{distinct}(id_1, \dots, id_f) \\ \text{distinct}(\text{methodname}(m_1), \dots, \text{methodname}(m_k)) \\ id \vdash m_i, i \in \{1, \dots, k\} \\ ??? \end{array}}{\vdash \text{class } id \text{ extends } id^P \{t_1 id_1; \dots; t_f id_f; m_1 \dots m_k\}}$$

Typechecking No-overloading in Minijava

$$\begin{array}{c} \text{distinct}(id_1, \dots, id_f) \\ \text{distinct}(\text{methodname}(m_1), \dots, \text{methodname}(m_k)) \\ id \vdash m_i, i \in \{1, \dots, k\} \\ \text{methodtype}(m_i, id^P) \neq \perp \Rightarrow \text{methodtype}(m_i, id^P) = \text{methodtype}(m_i, id) \\ \wedge \text{methodaccess}(m_i, id^P) = \text{methodaccess}(m_i, id), i \in \{1, \dots, k\} \\ \hline \vdash \text{class } id \text{ extends } id^P \{t_1 id_1; \dots; t_f id_f; m_1 \dots m_k\} \end{array}$$

$\text{methodtype}(D, id)$ returns

- the type of the method id in class D or one of its super-classes.
- $\perp$ if the method id is not defined in class D or one of its super-classes.

Soundness of a Type System

- A type system is *sound* if whenever $\vdash e : T$, then e will always evaluate to a value of type T in all executions.
- We want only sound rules, because then we can give a compile-time guarantee for the absence of *type errors*.
- Example of an unsound rule?

$$\frac{\vdash e_1 : int \quad \vdash e_2 : int}{\vdash e_1 + e_2 : int[]}$$

- A programming language is called *type safe* if a compiler of the language guarantees that compiled programs will run without type errors.
 - Can be either statically or dynamically typed.